


# x86 Paging - Complete Guide

## Virtual Address Structure (32 bits)



## Address Translation

Virtual 0xC0100500

Binary: 11000000 00010000 00000101 00000000

Bits 31-22: 11 0000 0000 = 768 (Directory Index)

Bits 21-12: 00 0100 0000 = 256 (Table Index)

Bits 11-0: 0101 0000 0000 = 0x500 (Offset)

Translation:

1. Look up `kernel_page_directory[768]` → Get address of page table
2. Look up `page_table[256]` → Get physical page address
3. Add offset: `physical_page + 0x500` = final physical address

## Page Directory

**Purpose:** Stores addresses of page tables

**Size:** 4KB (1024 entries × 4 bytes)

**Each entry covers:** 4MB of virtual address space

Index	Virtual Address Range	Points To
0	0x00000000 - 0x003FFFFF	Address of first_table
1	0x00400000 - 0x007FFFFF	Address of second_table (or NULL)
2	0x00800000 - 0x00BFFFFFF	NULL
...		
768	0xC0000000 - 0xC03FFFFF	Address of first_table (same as [0]!)
769	0xC0400000 - 0xC07FFFFF	Address of another table
...		
1023	0xFFC00000 - 0xFFFFFFFF	NULL

**Why index determines virtual range:** CPU extracts bits 31-22 from any virtual address to get the directory index.

## Page Table

**Purpose:** Stores physical addresses of actual memory pages

**Size:** 4KB (1024 entries  $\times$  4 bytes)

**Each entry covers:** 4KB of memory

Index	Physical Address
0	0x00000003 (phys: 0x00000000 + flags)
1	0x00001003 (phys: 0x00001000 + flags)
2	0x00002003 (phys: 0x00002000 + flags)
...	...
256	0x00100003 (phys: 0x00100000 + flags) ← Kernel here
...	...
1023	0x003FF003 (phys: 0x003FF000 + flags)

**Entry format:** Bottom 12 bits = flags, Top 20 bits = physical page address

## Example Setup: Identity Mapping

```
c

// Allocate page table (gets 4KB of physical memory)
uint32_t *first_table = (uint32_t*)pmm_alloc(1);

// Fill table: Map virtual page N → physical page N
for (int i = 0; i < 1024; i++) {
    uint32_t physical_addr = i * 0x1000; // Physical address
    first_table[i] = physical_addr | PAGE_PRESENT | PAGE_WRITE;
}

// Link table to page directory
kernel_page_directory[0] = ((uint32_t)first_table) | PAGE_PRESENT | PAGE_WRITE;
```

## What this does:

- Creates virtual addresses 0x00000000 - 0x003FFFFF (because we used directory[0])
- Maps them to physical addresses 0x00000000 - 0x003FFFFF
- Virtual = Physical (identity mapping)

## Higher-Half Kernel Mapping

c

```
// Map SAME table at index 768  
kernel_page_directory[768] = ((uint32_t)first_table) | PAGE_PRESENT | PAGE_WRITE;
```

**Result:** Now the SAME physical memory (0-4MB) is accessible via TWO virtual ranges:

- Virtual 0x00000000 - 0x003FFFFF (via directory[0])
- Virtual 0xC0000000 - 0xC03FFFFF (via directory[768])

Both point to physical 0x00000000 - 0x003FFFFF.

## Complete Example

Page Directory (at physical 0x9000)

Index	Value	What it does
0	0x00001003	Virtual 0x00000000- (first_table) 0x003FFFFF → Use   first_table for lookup
768	0x00001003	Virtual 0xC0000000- (first_table) 0xC03FFFFF → Use   first_table for lookup

Page Table: first\_table (at physical 0x1000)

Index	Value	Physical Address
0	0x00000003	0x00000000
1	0x00001003	0x00001000
256	0x00100003	0x00100000 (kernel!)
1023	0x003FF003	0x003FF000

## Translation Examples

**Virtual 0xC100500 → Physical ?**

Step 1: Extract bits

Directory Index = 768 (bits 31-22)

Table Index = 256 (bits 21-12)

Offset = 0x500 (bits 11-0)

Step 2: `kernel_page_directory[768] = 0x00001003`

→ `first_table` is at physical 0x00001000

Step 3: `first_table[256] = 0x00100003`

→ Physical page at 0x00100000

Step 4: Physical address =  $0x00100000 + 0x500 = 0x00100500$

## Virtual 0x00100500 → Physical ?

Step 1: Extract bits

Directory Index = 0

Table Index = 256

Offset = 0x500

Step 2: `kernel_page_directory[0] = 0x00001003`

→ `first_table` is at physical 0x00001000 (SAME TABLE!)

Step 3: `first_table[256] = 0x00100003`

→ Physical page at 0x00100000 (SAME RESULT!)

Step 4: Physical address =  $0x00100000 + 0x500 = 0x00100500$

**Both virtual addresses reach the same physical memory!**

## Key Points

1. **Page Directory entries:** Store addresses of page tables
2. **Page Table entries:** Store physical addresses of actual memory
3. **Virtual address bits 31-22:** Automatically select page directory index
4. **Virtual address bits 21-12:** Automatically select page table index
5. **You create virtual addresses:** By choosing which directory/table indices to fill
6. **CPU does translation:** You never manually convert - hardware does it automatically

## Memory Tracking

You need TWO separate tracking systems:

## Physical Memory Manager (PMM)

- Tracks which physical frames are free/used
- Returns physical addresses when you call `pmm_alloc()`

## Virtual Memory Manager (VMM)

- Tracks which virtual addresses are allocated
- Updates page tables when allocating virtual memory
- Uses bitmap or next-free-pointer to track virtual space

c

```
// Example allocation flow:  
void* vmm_alloc_pages(size_t count) {  
    // 1. Pick free virtual addresses (from your tracking)  
    uint32_t virt_addr = next_free_virtual_addr;  
  
    // 2. Get physical frames for each page  
    for (size_t i = 0; i < count; i++) {  
        uint32_t phys_frame = pmm_alloc(1); // Physical address  
        uint32_t virt_page = virt_addr + (i * 0x1000);  
  
        // 3. Map virtual → physical in page tables  
        map_page(virt_page, phys_frame, PAGE_PRESENT | PAGE_WRITE);  
    }  
  
    // 4. Update virtual address tracker  
    next_free_virtual_addr += count * 0x1000;  
  
    return (void*)virt_addr;  
}
```

## Enabling Paging

c

```
// 1. Set CR3 to page directory physical address
asm volatile("mov %0, %%cr3" :: "r"(kernel_page_directory));

// 2. Enable paging bit in CR0
uint32_t cr0;
asm volatile("mov %%cr0, %0" : "=r"(cr0));
cr0 |= 0x80000000; // Set PG bit
asm volatile("mov %0, %%cr0" :: "r"(cr0));
```

Now all memory accesses go through page tables!